

# Definitions


## What is ML ?

- **a definition:**

For a given **Task T**, a **machine** (algorithm) **A** obtains better **performance P** after an **experiment E**. (It has ***learned*** from it)  
(Experiment ~ data)

Yann LeCun's cake metaphor:

- **3 types** of learning :
  - **Supervised**
  - **Unsupervised**
  - **Reinforcement**  
(outside this course)



Reinforcement

Supervised

Unsupervised

# Today – Outline

- **Supervised Learning basics:**
  - Linear **regression**
  - Polynomial regression
- Lots of **Vocabulary**, notations
- Optimization basics: **Gradient Descent**
- **Supervised Learning**
  - Classification with the Perceptron (maybe)

# Today:

## Supervised Learning

Input:  $\vec{x}^{(n)} = (x_d^{(n)})_{d \in [1, \dots, D]}$ ,  $X = \{\vec{x}^{(n)}\}_{n \in [1, \dots, N]}$

- Expected Output:  $y^{GT}$  or  $t^{(n)}$  (**Ground Truth**)

Which kind of Task  $\rightarrow$  depends on  $t^{(n)}$

- **Model:**  $y^{predicted} \equiv \hat{y}^{(n)} = \sigma(f_{\Theta}(\vec{x}^{(n)}))$

fct.  $f_{\Theta}$  is **parameterized** by **parameters**

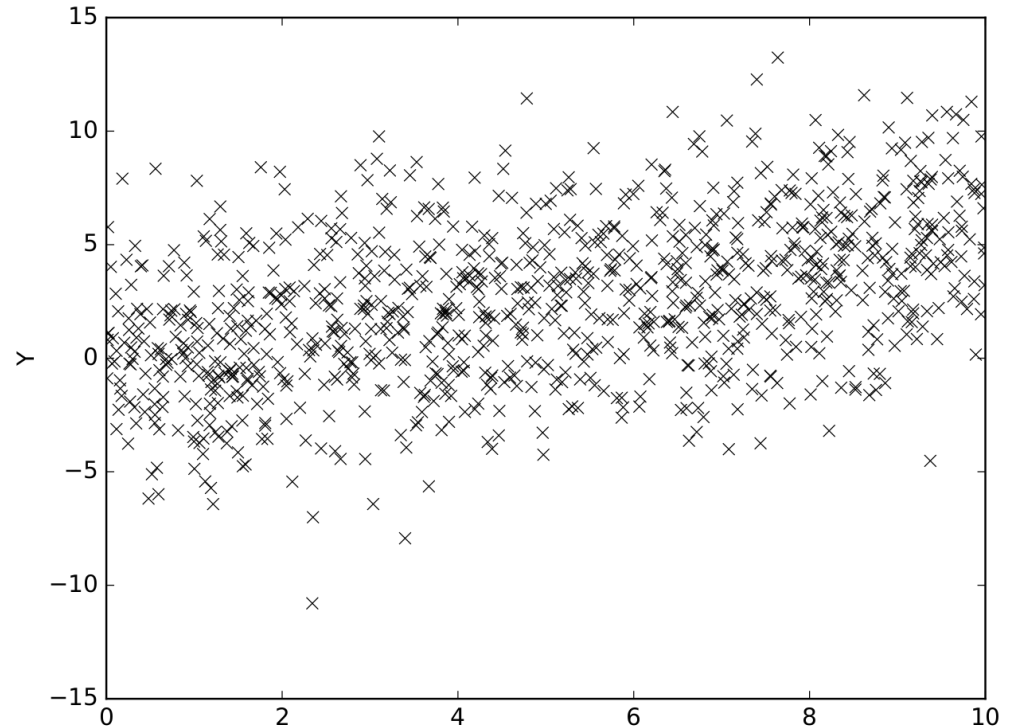
- **Learning** : finding *optimal* parameters to minimize discrepancy between  $\hat{y}$  and Ground Truth  $t$

$$\Theta^* = \operatorname{argmin}_{\Theta} \left( \sum_n^N \mathcal{L}(\hat{y}_n, t_n) \right)$$

- **Cost Function (loss function)** : to be chosen <sub>3</sub>

# Supervised Learning: Regression

Pairs of data  
points  $\vec{x}^{(n)} = (x_1^{(n)}, x_2^{(n)})$



→ Relationship  $f(x)=y$  ?

→ **Regression**

- **linear:**  $f_{a,b}(x) = ax + b$  or  $f_{\vec{a},b}(\vec{x}) = \vec{a} \cdot \vec{x} + b$

- **polynomial:**

(degree  $P$ )

$$f_{\Theta}(\vec{x}) = \sum_{p=0}^P \vec{\theta}_p \cdot \vec{x}^p$$

# More Vocabulary

(+case of Regression)

Input:  $\vec{x}^{(n)} = (x_d^{(n)})_{d \in [1, \dots, D]}, X = \{\vec{x}^{(n)}\}_{n \in [1, \dots, N]}$

- *Ground Truth:*  $t^{(n)} \in \mathbb{R}, T = \{t^{(n)}\}_{n \in [1, \dots, N]}$

**Continuous** output  $\rightarrow$  Task is **Regression**

- **Model:** e.g. a polynomial function of the input :  $f_{\Theta}(\vec{x}) = \sum_{p=0}^P \vec{\theta}_p \cdot \vec{x}^p$

- Parameters:  $\Theta = \{\theta_0, \theta_{d,p} / d = 1, \dots, D, p = 1, \dots, P\}$

- **Prediction:** simply  $\hat{y}_n = f_{\Theta}(\vec{x}_n)$



- Learning **Algorithm:**

$$Card(\Theta) = D.P + 1$$

- **Initialization:**  $\Theta = \Theta_0$

- Minimize some Loss  $\mathcal{L}(\hat{y}_n, t_n)$  (to choose)

- For this, use some minimization scheme (Grad. Desc.)

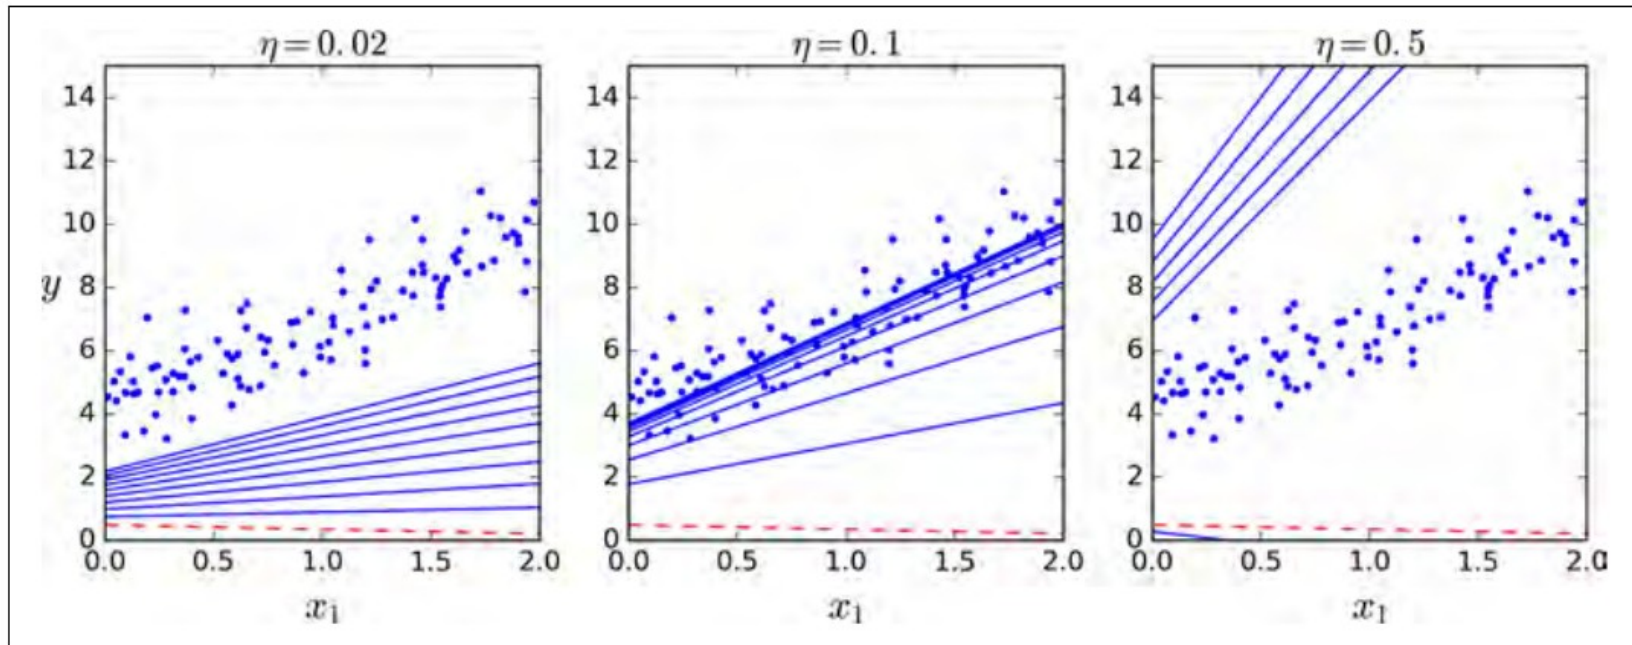
# Supervised Learning: Regression

- We can choose: **Least Squares**

Single data point Loss:  $\mathcal{L}(f_{\Theta}(\vec{x}_n), t_n) = (\vec{f}(\vec{x}^{(n)}) - t^{(n)})^2$

$$\text{Global Loss: } \mathcal{L}(X, T) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}(f_{\Theta}(\vec{x}_n), t_n)$$

- **Gradient Descent:**









# Gradient Descent

## short reminder


- I have a function  $J(\theta)$  and want to find the value  $\theta^*$  for which  $J(\theta)$  is minimum

# What is the gradient ?

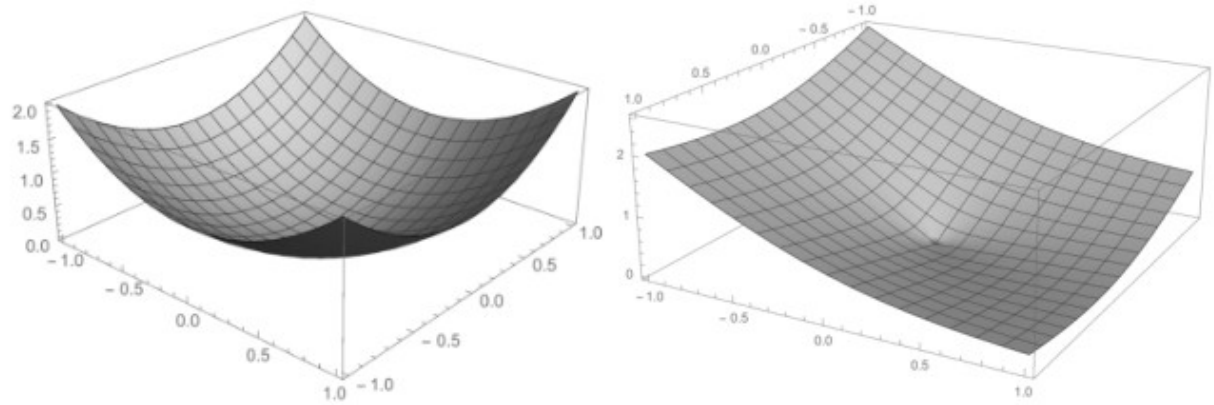





$x^3 + y^3$ , et  $-(x^2 + y^2) + (x^4 + y^4)$  :



$x^2 + y^2$  et  $||\vec{x}|| - a\vec{w} \cdot \vec{x} = (x^2 + y^2)^{1/2} - aw_1x - aw_2y$ , avec  $a = 3, w_1 = 0.1, w_2 = 0.3$  :



$e^{-x^2}y^2$ , et  $-(x^2 + y^2) + (x^2 + y^2)^2$



# Gradient Descent

- Limitations:
  - at best, converges to *one of the **local** minima*
  - typically converges to the minimum of the local basin of attraction we are in
  - there may be many local minima. The best one may not be close to our (random) starting position...
  - it may never converge (diverge or continuously go down)
  - it goes in the steepest direction (from the local point) → is also called “*steepest descent*”

# Least Squares

$$LSE = \frac{1}{N} \sum_{n=1}^N (\vec{f}_{\Theta}(x^{(n)}) - \vec{y}^{(n)})^2, \text{ with } f_{\Theta}(\vec{x}) = \sum_{p=0}^P \vec{\theta}_p \cdot \vec{x}^p$$

Case  $P = 1$

# Trick: Augmented data

- Add 1's into  $X$  to take care of the offset, once and for  
→ get cleaner equations (and cleaner code) !

# References:

## **Linear regression (by G.D.)**

→ Bishop book, page 143-144, section 3.1.3  
(sequential learning)

→ [https://en.wikipedia.org/wiki/Least\\_squares#Linear\\_least\\_squares](https://en.wikipedia.org/wiki/Least_squares#Linear_least_squares)

- **Gradient Descent (assumed known)**

→ catch up lesson:


[https://en.wikipedia.org/wiki/Gradient\\_descent](https://en.wikipedia.org/wiki/Gradient_descent)

# Pause-Questions

# First nuance: various **optimization strategies**

- **Examples** seen one by one:  
*“Online” learning*
- Examples seen all at once  $(m=N)$   
**globale update (optimization viewpoint)**
- Intermediate solution: **batch size**  $m$ ,  $(m>1)$   
**mini-batch Gradient Descent**
- **Stochastic Gradient Descent (SGD)**:  $(m=1)$   
~looks like *Online* (but more random)

# SGD vs *mini-batch* vs *full batch*





# Key concepts

- **Supervised Learning**
- **Regression**
- **Task, Model**, parameters, **prediction**/decision, input **feature**
- [SGD, mini-batch, full batch, Online]