

NoSQL

Chapter contents:

- NoSQL: General principles
- NoSQL data models: JSON (and XML)
- Data exchange and serialization

Table of contents

NoSQL

- NoSQL: General principles
- NoSQL data models: JSON (and XML)
- Data exchange and serialization

Latency numbers every programmer should know (J.Dean)

Latency numbers (2012)

L1 cache reference	0.5 ns	
Branch mispredict	5 ns	
L2 cache reference	7 ns	
Mutex lock/unlock	25 ns	
Main memory reference	100 ns	
Compress 1K bytes with Zippy	3,000 ns	= 3 μ s
Send 2K bytes over 1 Gbps network	20,000 ns	= 20 μ s
SSD random read	150,000 ns	= 150 μ s
Read 1 MB sequentially from memory	250,000 ns	= 250 μ s
Round trip within same datacenter	500,000 ns	= 0.5 ms
Read 1 MB sequentially from SSD*	1,000,000 ns	= 1 ms
Disk seek	10,000,000 ns	= 10 ms
Read 1 MB sequentially from disk	20,000,000 ns	= 20 ms
Send packet CA->Netherlands->CA	150,000,000 ns	= 150 ms


2019: for a 4kB block:

Latency numbers (2019)

Fast NVMe (Optane)	7 μ s
Fast NVMe (Z-SSD)	12 μ s
Round trip TCP packet on 10Gb Ethernet ...	20-50 μ s
NVMe Flash SSD	80 μ s

Throughput numbers on key-value stores

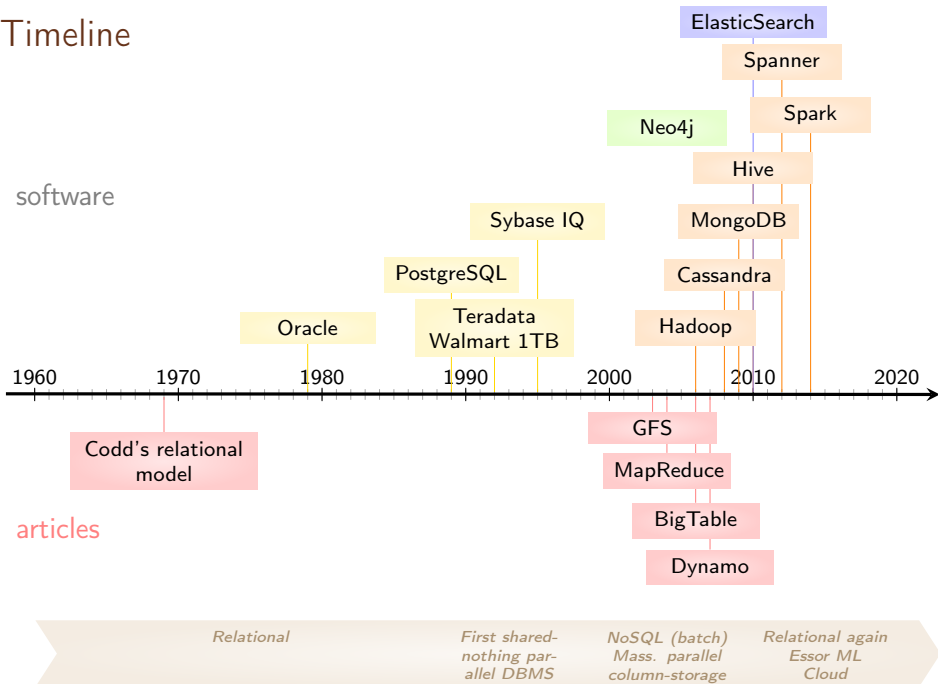
Achieved read throughput
on a 20-core 24-device system



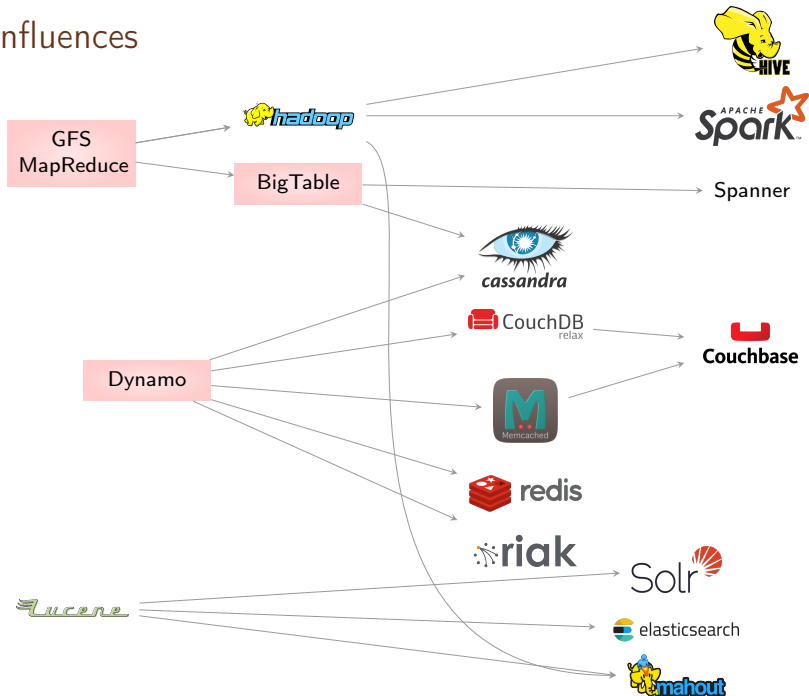
[<https://www.usenix.org/system/files/fast19-kourtis.pdf>]

Timeline



software



Influences



Google trends



Google Trends: search comparison

Distributed Databases

Why *distribute*?

- parallelism (=performance)
- scalability
- availability: accessibility and fault tolerance (cloud)
- optimize for different hardware, distribute geographically,...

How to distribute?

- **sharding** \simeq horizontal partitioning
- but also **replication**

Implementation challenges:

- decentralized architecture maintain coherence between copies, task and data partitioning
- **Shared nothing architecture** (nots *shared disk*, not *shared memory pool*).
how to chose the partitioning

Types of parallelism

Parallelism in DBMS has a long history:

- **inter-operator:** every CPU computes a query operation (pipeline).
Volcano model – a query operation sends the output directly to the next operation.
- **intra-operator:** every CPU computes the entire query on a part of the data
- **inter-query:** several queries executed in parallel

Since then:

- large scale data – distributed computing on a large amount of computers.
- *Shared nothing architecture* (neither *shared disk* nor *shared memory pool*).

Replication

Objectives: reliability, read performance.

Techniques:

- RAM+logs on disk: write-ahead logs (WAL)
- generally, asynchronous (eventual consistency)
- sometimes synchronized (but can have slow updates)
- versioning (vector clocks)
- network state (faults,...): gossip
- fault recovery: consensus (Paxos)

Ex: MongoDB: asynchronous, WAL.

In distributed DBMS:

PostgreSQL (WAL), MariaDB, Oracle (materialized views), SQL Server...

Often admin level choices (number of masters, synchronization,...).

Challenges when distributing

- good data partition/replication
- coherence (trade-off between performance and integrity when dealing with reads and writes)
- distributing computation tasks (to minimize data exchange)
- fault tolerance
- transaction control
- data privacy

Distributed architectures

Master-slave

- MongoDB: server mongos/mongod
- HDFS: NameNode/DataNodes
- BigTable

Without master servers

- Dynamo
- Cassandra
- (BitTorrent)

Partitioning: *how to distribute data ?*

↪ technique in DynamoDB: coherent hashing.

Consistency, Availability, Partition tolerance

Ideally, distributed database systems would need to provide 3 guarantees:

Consistency:

Every read request to the system receives data corresponding to the most recent read.

Availability:

The system must answer any query to the system, even if the answer is wrong or outdated.

Partition tolerance:

The system must answer queries even under arbitrary failures of distributed nodes or of messages between nodes.


Consistency, Availability, Partition tolerance

CAP Theorem (Consistency, Availability, Partition tolerance):

A system cannot have all 3: when data is partitioned, we cannot guarantee both availability and consistency.

- In a distributed setting, ACID (Atomic, Consistence, Isolation, Durable) guarantees consistency while sacrificing availability.
- NoSQL DBs use rather the BASE consistency model: (Basic Availability, Soft state, Eventual consistency):
 - Basic availability: data is always available,
 - Soft-state: different copies are not always consistent
 - Eventual consistency: after a while (if no changes), the system is consistent

In other words: the system has a mechanism to reconcile (sooner or later) all the versions.



NoSQL

NoSQL: Not SQL or (more often) Not Only SQL

Data model not fixed (and not relational) – e.g., *key/value pairs*,
value: document or hashmap, ...

Rough taxonomy (not standardized):

- Key-Value: (Redis, Memcached, Riak)
- Document: (MongoDB, CouchDB)
- Column: *Column-family*: Cassandra,
Column-oriented: SAP Hana, MonetDB
- Graph: (Neo4j)

Principal characteristic: very vague classification, under constant evolution.

NoSQL (2)

- query language is no longer only SQL:
method calls in programming languages (object+functional).
- new software stacks:
 - Web development:
LAMP (Linux, Apache, Mysql, PHP) ¹
MEAN (MongoDB, ExpressJS, AngularJS, NodeJS)
 - and for Big data? Not settled yet
e.g, SMACK (Spark, Apache Mesos, Akka, Cassandra, and Apache Kafka)
- NoSQL means also no standard!

¹+variants: other OS (Windows), server (Nginx), storage (MariaDB), script

NoSQL design principles

NoSQL is focused on data exchange and distribution.

We want *autonomous* data:

- denormalization
- no joins (autonomous documents)

Typical application: machine learning training data, graph data.

Typical NoSQL database: a “bunch” of documents

E.g., scientific papers: each article contains the entire information (authors, etc.)

Comparing NoSQL to relational databases

Roughly:

- ✓ easy to distribute migrate
- ✓ performance-oriented: scalability (>TB), real-time
- ✓ easy to design (usually, no modeling needed)
- ✓ high availability
- ✓ somewhat un-structured data (multimedia, graphs), for which relational DBs are not adapted
- ✓ very useful if many reads, few updates
- ✓ deals directly with programming APIs

- ✗ no standard query language: need to program
- ✗ asymmetric: some data access patterns are favoured
- ✗ no (or fixed) schema
- ✗ no transactions

Table of contents

NoSQL

- NoSQL: General principles
- NoSQL data models: JSON (and XML)
- Data exchange and serialization

Relational data model, XML, JSON

Heavily inspired by <http://b3d.bdpedia.fr>

Data format for structured data: XML and JSON.

Made principally for data exchange.

XML: eXtensible Markup Language

JSON: JavaScript Object Notation

XML

```
<personne>
  <nom>Dupond</nom>
  <tels>
    <tel>0612304056</tel>
    <tel>0269159002</tel>
  </tels>
</personne>
```

JSON

```
{
  "nom": "Dupond",
  "tels": [0612304056, 0269159002]
}
```

✓ very rich ecosystem:
XLST, XQuery, SVG, RSS

✓ simple, compact

✗ limited (in terms of types)

✗ verbose

✗ complex

JSON

JSON object: set of *key value pairs*

keys are strings

values can be:

- JSON object
 - array of values
 - string
 - number
 - boolean
 - null
- } atomic types

Key value pairs exist in different versions:

JSON objects, XML elements, associative arrays (PHP), hash map (Java), dictionaries (Python)...

JSON validation

JSON

```
{
  "menu": {
    "id": "file",
    "value": "File",
    "popup": {
      "menuitem": [
        {
          "value": "New",
          "onclick": "CreateNewDoc()"
        },
        {
          "value": "Open",
          "onclick": "OpenDoc()"
        },
        {
          "value": "Close",
          "onclick": "CloseDoc()"
        }
      ]
    }
  },
  "mixed_list": ["aabb", 2018, true, [1,2,3]]
}
```

Multiple validators available (also for XML, HTML, etc).

E.g.: <http://jsonlint.com>

JSON Schema

JSON

```
{
  "checked": true,
  "dimensions": {
    "width": 5,
    "height": [1,2]
  }
}
```

JSON Schema

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "properties": {
    "checked": {
      "$id": "/properties/checked",
      "type": "boolean",
      "title": "The Checked Schema ",
      "default": false,
      "examples": [ true ]
    },
    ...
  }
}
```

JSON Schema describes (using JSON) the structure of a JSON object (like XSD for XML).

Examples:

- schema validator: <https://www.jsonschemavalidator.net>
- automatic schema inference: <https://www.jsonschema.net>

Table of contents

NoSQL

- NoSQL: General principles
- NoSQL data models: JSON (and XML)
- Data exchange and serialization

(De-)serializing JSON data (Javascript)

Sending data:

```
//JSON.stringify(v) converts value v into JSON string:  
var myObj = { "nom": "Dupond", "age": 31, "ville": "Paris" };  
var myJSON = JSON.stringify(myObj);  
window.location = "demo_json.php?x=" + myJSON;
```

```
JSON.stringify(new Date(2006, 0, 2, 15, 4, 5))  
// ""2006-01-02T15:04:05.000Z""
```

Receiving data:

```
//JSON.parse(s) decodes string s (that represents JSON object)  
// into JavaScript object  
var myJSON = '{ "nom": "Dupond", "age": 31, "ville": "Paris" }';  
var myObj = JSON.parse(myJSON);  
document.getElementById("demo").innerHTML = myObj.nom;  
// "Dupond"
```

(De-)serializing JSON data (Python)

For simple objects (dict, list...):

```
import json

# converts a dict into JSON string:
mydict = {'nom': 'Dupond', 'age': 31, 'ville': 'Paris'}
myJson = json.dumps(data)

# and conversely to recreate object:
json.loads(encoded_hand)
```

```
import json

# serializes a dict into JSON file:
with open("data_file.json", "w") as f:
    json.dump(data, f)

# and parses JSON file into dict:
with open("data_file.json", "r") as f:
    data = json.load(f)
```

(De-)serializing JSON data (Python)

For a class:

```
import json
""" solution 1: define an encoder for the class """
data = maclasse(...)
from json import JSONEncoder
class maclasseEncoder(JSONEncoder):
    def default(self, o):
        return o.__dict__ # not really robust

myJson = json.dumps(data, cls=maclasseEncoder)
# or myJSON = maclasseEncoder.encode(data)

""" solution2: make the class serializable by implementing the toJson method """
def toJson(self):
    return json.dumps(self, default=lambda o: o.__dict__)
myJson = json.dumps(data.toJson())

""" solution 3: use module 'jsonpickle' """
```

Data serialization framework (RPC)

Protocol Buffers (proto2, proto3):

- messages are key-value pairs
- define, in a file .proto, the message structure
- the message is then compiled in a programming language of choice
- compact format
- used by gRPC (gRPC+protobuf faster than Rest+JSON)

Thrift: then

- we define a message structure then we compile it into an object
- more languages available than Protocol Buffers (?)

Avro

- schema defined in JSON, dynamic, not compiled

Data exchange using REST APIs

Representational state transfer

Restful API contains:

- a resource description URI
- HTTP methods
- MIME types: JSON, XML, but also data structure (pages, sorting, ...) – can take considerable implementation effort

RPC API vs JSON HTTP API

RPC use cases:

- ✓ micro-services (RPC is low latency, high debit, so low network volume)
- ✓ streaming (real-time)

Disadvantages:

- ✗ cannot call services directly via HTTP (=browser)
- ✗ binary format, so not human readable

[<https://docs.microsoft.com/fr-fr/aspnet/core/grpc/comparison>]

Using a message broker

Publish & subscribe platforms

Intermediary between providers and receivers (subscribers)

- ✓ exchange is simplified (providers and receivers are independent)
- ✓ allows asynchronous communication



Compared to RPC:

- ✗ no direct communication
- ✗ sometimes needs heavy architecture (Kafka vs. gRPC)

message queue: each message read once, or *pub/sub* – only subscribers are read

Example of a message agent

- data streams
- guarantees receiving the message in the same order
- partitioning and replication
- objective: low latency, high debit

Uses a batch messaging principle via TCP

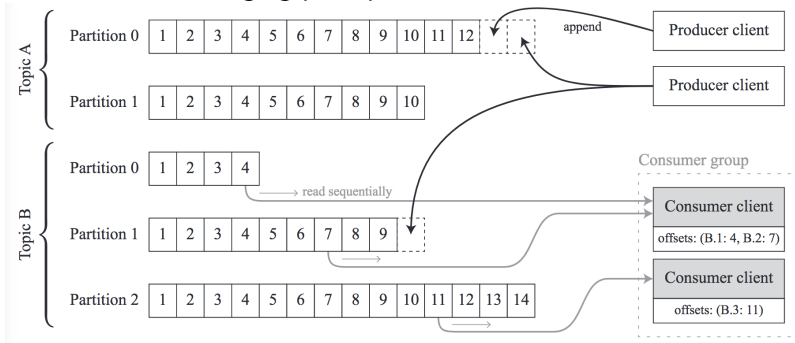


Figure 1: A Kafka *topic* is divided into *partitions*, and each partition is a totally ordered sequence of *messages*.

<https://www.repository.cam.ac.uk/bitstream/handle/1810/286031/kafka-debull15.pdf?sequence=1>

Column stores

4 example projects  for serializing column data:

Arrow  : *in-memory*

- allows direct access to column data, without the need to access row-by-row
- vectorization approach, random access, zero-copy
- Feather: adaptation of Arrow for file storage

Parquet  **Parquet** : *on disk*

- uses compression approaches


Kudu  : *on disk (+ cache)*

- optimised for updates

ORC 

- used to optimise Hive and MapReduce
- compression, indexes, predicate pushdown

Google trends



Apache Avro started in 2009 !

Bibliography

- Semi-structured data

https://www.w3schools.com/js/js_json_intro.asp

An overview of JSON data model and the typical operations on JSON:

JSON: data model, query languages and schema specification, Bourhis et al., PODS 2017

<https://arxiv.org/pdf/1701.02221.pdf>

<http://b3d.bdpedia.fr/docstruct.html>

An experimental comparison of serialization mechanisms:

<http://labs.criteo.com/2017/05/serialization/>

Course slides on Protocol Buffers, Thrift, Avro:

https://ganges.usc.edu/pgroupW/images/a/a9/Serializarion_Framework.pdf

Arrow vs Parquet: <http://wesmckinney.com/blog/arrow-columnar-abadi/>

also here: <http://dbmsmusings.blogspot.fr/2018/03/>

[an-analysis-of-strengths-and-weaknesses.html](http://dbmsmusings.blogspot.fr/2018/03/an-analysis-of-strengths-and-weaknesses.html)